

Totpal - Project 2 Code

February 15, 2026

Title: ICU Occupancy Change Prediction in California

Author: Allyson Totpal

Date: 15 February 2026

Modified By: Allyson Totpal

Description: Forecasting medium-term hospital capacity strain in California using respiratory emergency department surveillance signals.

```
[2]: import pandas as pd
import numpy as np

# read in data
hhs = pd.read_csv('/Users/smooshii/DSC680/
    ↪COVID-19_Reported_Patient_Impact_and_Hospital_Capacity_by_Facility_20260208.
    ↪csv')
nssp = pd.read_csv('/Users/smooshii/DSC680/
    ↪NSSP_Emergency_Department_Visit_Trajectories_by_State_and_Sub_State_Regions-_COVID-19,_Flu,
    ↪csv')
cdc_dash = pd.read_csv('/Users/smooshii/DSC680/respiratory-virus-dashboard.csv')

/var/folders/_p/j05_fthn2yqfx9lkpltf1p0h0000gn/T/ipykernel_75917/1197078806.py:5
: DtypeWarning: Columns (0,3) have mixed types. Specify dtype option on import
or set low_memory=False.
    hhs = pd.read_csv('/Users/smooshii/DSC680/COVID-
19_Reported_Patient_Impact_and_Hospital_Capacity_by_Facility_20260208.csv')
```

```
[3]: nssp.head()
```

```
[3]:      week_end geography     county percent_visits_combined \
0  2022-10-01   Alabama      Bibb            NaN
1  2022-10-01   Alabama    Calhoun            NaN
2  2022-10-01   Alabama   Chilton            NaN
3  2022-10-01   Alabama  Cleburne            NaN
4  2022-10-01   Alabama    Coosa            NaN

  percent_visits_covid  percent_visits_influenza  percent_visits_rsv \

```

```

0           NaN           NaN           NaN
1           NaN           NaN           NaN
2           NaN           NaN           NaN
3           NaN           NaN           NaN
4           NaN           NaN           NaN

percent_visits_smoothed_combined percent_visits_smoothed_covid \
0           NaN           NaN
1           NaN           NaN
2           NaN           NaN
3           NaN           NaN
4           NaN           NaN

percent_visits_smoothed_influenza percent_visits_smoothed_rsv \
0           NaN           NaN
1           NaN           NaN
2           NaN           NaN
3           NaN           NaN
4           NaN           NaN

ed_trends_covid ed_trends_influenza     ed_trends_rsv \
0 Data Unavailable   Data Unavailable   Data Unavailable
1 Data Unavailable   Data Unavailable   Data Unavailable
2 Data Unavailable   Data Unavailable   Data Unavailable
3 Data Unavailable   Data Unavailable   Data Unavailable
4 Data Unavailable   Data Unavailable   Data Unavailable

hsa \
0 Jefferson (Birmingham), AL - Shelby, AL
1 Calhoun (Anniston), AL - Cleburne, AL
2 Jefferson (Birmingham), AL - Shelby, AL
3 Calhoun (Anniston), AL - Cleburne, AL
4 Talladega, AL - Clay, AL

hsa_counties hsa_nci_id    fips \
0 Bibb, Blount, Chilton, Cullman, Jefferson, She...      150 1,007
1                           Calhoun, Cleburne            177 1,015
2 Bibb, Blount, Chilton, Cullman, Jefferson, She...      150 1,021
3                           Calhoun, Cleburne            177 1,029
4                           Clay, Coosa, Talladega        241 1,037

trend_source BuildNumber
0      HSA 2026-02-06
1      HSA 2026-02-06
2      HSA 2026-02-06
3      HSA 2026-02-06
4      HSA 2026-02-06

```

0.0.1 NSSP Preprocessing

```
[5]: # standardize column names
nssp.columns = [c.strip() for c in nssp.columns]
nssp.head()

# parse dates
nssp['week_end'] = pd.to_datetime(nssp['week_end'], errors = 'coerce')

# filter to CA only
nssp = nssp[nssp['geography'].str.lower() == 'california'].copy()

# drop high granularity geography
nssp = nssp.drop(columns = ['county', 'hsa'], errors = 'ignore')

[6]: # candidate Ed signal columns
ed_cols = ['percent_visits_influenza', 'percent_visits_rsv', \
    ↴'percent_visits_combined', \
        'percent_visits_smoothed_influenza', 'percent_visits_smoothed_rsv', \
            ↴'percent_visits_smoothed_combined']
keep_cols = ['week_end'] + [c for c in ed_cols if c in nssp.columns]
nssp = nssp[keep_cols]

# convert to numeric
for c in ed_cols:
    if c in nssp.columns:
        nssp[c] = pd.to_numeric(nssp[c], errors = 'coerce')

[7]: # aggregate to one row per week_end
nssp_week = (nssp.groupby('week_end', as_index = False)
    .mean()
    .sort_values('week_end')
)
nssp_week.rename(columns = {'week_end': 'week_ending'}, inplace = True)
nssp_week.head()

[7]:   week_ending  percent_visits_influenza  percent_visits_rsv  \
0  2022-10-01              0.21              0.13
1  2022-10-08              0.27              0.23
2  2022-10-15              0.43              0.39
3  2022-10-22              0.65              0.60
4  2022-10-29              1.01              0.80

    percent_visits_combined  percent_visits_smoothed_influenza  \
0                  1.51                      0.15
1                  1.58                      0.21
2                  1.83                      0.31
```

3	2.31	0.45
4	2.95	0.70
percent_visits_smoothed_rsv	percent_visits_smoothed_combined	
0	0.10	1.55
1	0.16	1.55
2	0.26	1.65
3	0.41	1.91
4	0.59	2.36

```
[8]: # engineer rolling means + week-over-week deltas
for c in ed_cols:
    if c in nssp_week.columns:
        nssp_week[f'{c}_roll2'] = nssp_week[c].rolling(2, min_periods = 1).mean()
        nssp_week[f'{c}_roll3'] = nssp_week[c].rolling(3, min_periods = 1).mean()
        nssp_week[f'{c}_wow_delta'] = nssp_week[c].diff()
```

0.0.2 CDC Dashboard Preprocessing

```
[10]: cdc_dash['WEEKENDING'] = pd.to_datetime(cdc_dash['WEEKENDING'], errors = 'coerce')

# identify useful columns
cols_of_interest = ['WEEKENDING', 'FLU_ED_VISITS', 'RSV_ED_VISITS', 'FLU_ADMITTISSONS', 'RSV_ADMITTISSONS', 'POP']
cdc_dash = cdc_dash[[c for c in cols_of_interest if c in cdc_dash.columns]]
```

```
[11]: # convert to numeric
for c in cdc_dash.columns:
    if c != 'WEEKENDING':
        cdc_dash[c] = pd.to_numeric(cdc_dash[c], errors = 'coerce')
```

```
[12]: # aggregate to state-week
cdc_dash_week = (cdc_dash.groupby('WEEKENDING', as_index = False)
                 .sum(numeric_only = True)
                 .sort_values('WEEKENDING')
                 )
cdc_dash_week.rename(columns = {'WEEKENDING': 'week_heading'}, inplace = True)
cdc_dash_week.head()
```

```
[12]: week_heading FLU_ED_VISITS RSV_ED_VISITS POP
0 2023-07-08 0.1 0.0 0.0
1 2023-07-15 0.1 0.0 0.0
2 2023-07-22 0.1 0.0 0.0
```

```

3 2023-07-29          0.1          0.0  0.0
4 2023-08-05          0.1          0.0  0.0

```

```
[13]: # compute per-100k rates if POP exists
if 'POP' in cdc_dash_week.columns:
    pop = cdc_dash_week['POP'].replace(0, np.nan)

    for c in ['FLU_ED_VISITS', 'RSV_ED_VISITS', 'FLU_ADMISSESIONS', ↴
    'RSV_ADMISSESIONS']:
        if c in cdc_dash_week.columns:
            cdc_dash_week[f'{c}_per_100k'] = (cdc_dash_week[c] / pop) * 100_000
```

```
[14]: # rolling features on rate columns
rate_cols = [c for c in cdc_dash_week.columns if c.endswith('_per_100k')]

for c in rate_cols:
    cdc_dash_week[f'{c}_roll2'] = cdc_dash_week[c].rolling(2, min_periods = 1).mean()
    cdc_dash_week[f'{c}_roll3'] = cdc_dash_week[c].rolling(3, min_periods = 1).mean()
    cdc_dash_week[f'{c}_wow_delta'] = cdc_dash_week[c].diff()
```

0.0.3 HHS Preprocessing

```
[16]: # select columns needed
hhs_cols = ['collection_week', 'state', ↴
    'all_adult_hospital_inpatient_beds_7_day_sum',
    'inpatient_beds_used_7_day_sum', ↴
    'staffed_adult_icu_bed_occupancy_7_day_avg']
hhs = hhs[hhs_cols].copy()
```

```
[17]: # parse week & filter to CA
hhs.loc[:, 'collection_week'] = pd.to_datetime(hhs.loc[:, 'collection_week'], ↴
    errors = 'coerce')
hhs = hhs[hhs['state'] == 'CA'].copy()

# columns to clean
count_cols = ['all_adult_hospital_inpatient_beds_7_day_sum', ↴
    'inpatient_beds_used_7_day_sum']
raw_rate_cols = 'staffed_adult_icu_bed_occupancy_7_day_avg'
```

```
[18]: def _clean_numeric_series(s: pd.Series) -> pd.Series:
    '''robust numeric cleaning'''
    return pd.to_numeric(
        s.astype(str)
        .str.replace(',', '', regex = False)
        .str.replace('%', '', regex = False)
```

```

    .str.strip()
    .replace({'': np.nan, 'nan': np.nan, 'None': np.nan}),
    errors = 'coerce'
)

```

[19]:

```

# clean numeric fields
for c in count_cols + [raw_rate_cols]:
    hhs[c] = _clean_numeric_series(hhs[c])

# aggregate to week
hhs_week = (hhs.groupby('collection_week', as_index = False)
            .agg({**{c: 'sum' for c in count_cols},
                  raw_rate_cols: 'mean'})
            .sort_values('collection_week')
            .rename(columns = {'collection_week': 'week_ending'}))
)

```

[20]:

```

hhs_week['icu_occ_raw'] = hhs_week[raw_rate_cols].abs() / 1_000_000

# winsorize ICU occupancy at p99
cap_99 = hhs_week['icu_occ_raw'].quantile(0.99)
hhs_week['icu_occ_wins'] = hhs_week['icu_occ_raw'].clip(upper = cap_99)

# final ICU occupancy feature
hhs_week['icu_occ'] = hhs_week['icu_occ_wins']

# inpatient occupancy ratio
hhs_week['inpatient_occ'] = (hhs_week['inpatient_beds_used_7_day_sum'] /
                             □
                             ↵hhs_week['all_adult_hospital_inpatient_beds_7_day_sum'].replace(0, np.nan))

# high strain flag
p95 = hhs_week['icu_occ_wins'].quantile(0.95)
hhs_week['high_strain'] = (hhs_week['icu_occ_wins'] >= p95).astype(int)

# create standardized Monday-start week key for joining
hhs_week['week'] = hhs_week['week_ending'].dt.to_period('W-MON').dt.start_time

```

[21]:

```

print('Weeks:', len(hhs_week))
print('Non-missing icu_occ:', int(hhs_week['icu_occ'].notna().sum()))
print('ICU occ wins cap (p99):', cap_99)
print(hhs_week['icu_occ'].describe())
print('p95 threshold:', p95)
print(hhs_week['high_strain'].value_counts())

hhs_week.head()

```

Weeks: 218

```

Non-missing icu_occ: 198
ICU occ wins cap (p99): 0.1930038756231166
count      198.000000
mean       0.140260
std        0.029996
min        0.059637
25%        0.121798
50%        0.142028
75%        0.163716
max        0.193004
Name: icu_occ, dtype: float64
p95 threshold: 0.18143368775902988
high_strain
0      208
1      10
Name: count, dtype: int64

```

```
[21]: week_ending  all_adult_hospital_inpatient_beds_7_day_sum \
0  2020-02-02          0.0
1  2020-03-01          0.0
2  2020-03-08          0.0
3  2020-03-15          0.0
4  2020-03-22          0.0

inpatient_beds_used_7_day_sum  staffed_adult_icu_bed_occupancy_7_day_avg \
0                      77.0           NaN
1                     50.0           NaN
2                     594.0          NaN
3                    2127.0          NaN
4                   5973.0          NaN

icu_occ_raw  icu_occ_wins  icu_occ  inpatient_occ  high_strain    week
0         NaN        NaN     NaN        NaN        NaN  0 2020-01-28
1         NaN        NaN     NaN        NaN        NaN  0 2020-02-25
2         NaN        NaN     NaN        NaN        NaN  0 2020-03-03
3         NaN        NaN     NaN        NaN        NaN  0 2020-03-10
4         NaN        NaN     NaN        NaN        NaN  0 2020-03-17

```

0.0.4 Merge Weekly Datasets

```
[23]: # merge on week_ending
hhs_week['week'] = hhs_week['week'].dt.to_period('W-MON').dt.start_time
nssp_week['week'] = nssp_week['week'].dt.to_period('W-MON').dt.start_time
cdc_dash_week['week'] = cdc_dash_week['week'].dt.to_period('W-MON').dt.
    ↪start_time
```

```
[24]: data = (hhs_week
             .merge(nssp_week, on = 'week', how = 'left')
             .sort_values('week')
             .reset_index(drop = True)
         )
```

0.0.5 Build Forecasting and Modeling

```
[26]: # create forecast target (3 weeks ahead)
forecast_horizon = 3
data['icu_occ_t_plus_3w'] = data['icu_occ'].shift(-forecast_horizon)

# drop rows where future target is missing
model_df = data.dropna(subset = ['icu_occ_t_plus_3w']).copy()
print('Model Rows: ', len(model_df))
```

Model Rows: 198

```
[27]: # select predictors
feature_candidates = ['icu_occ', 'FLU_ED_VISITS_per_100k',
                      'RSV_ED_VISITS_per_100k',
                      'FLU_ED_VISITS_per_100k_roll2',
                      'FLU_ED_VISITS_per_100k_roll3',
                      'RSV_ED_VISITS_per_100k_roll2',
                      'RSV_ED_VISITS_per_100k_roll3']
features = [f for f in feature_candidates if f in model_df.columns]
print('Using features: ', features)
```

Using features: ['icu_occ']

```
[28]: # build x and y
x = model_df[features].copy()
y = model_df['icu_occ_t_plus_3w']

# drop rows with missing predictors
mask = x.notna().all(axis = 1)
x = x[mask]
y = y[mask]

print('Final Modeling Rows: ', len(x))
```

Final Modeling Rows: 195

```
[29]: # train/test split
split_idx = int(len(x) * 0.8)

x_train = x.iloc[:split_idx]
x_test = x.iloc[split_idx:]
```

```
y_train = y.iloc[:split_idx]
y_test = y.iloc[split_idx:]
```

```
[30]: from sklearn.linear_model import Ridge
from sklearn.preprocessing import StandardScaler
from sklearn.pipeline import Pipeline
from sklearn.metrics import mean_squared_error, mean_absolute_error

# ridge regression (standardized)
ridge_model = Pipeline([('scaler', StandardScaler()),
                        ('ridge', Ridge(alpha = 1.0))])
ridge_model.fit(x_train, y_train)
ridge_preds = ridge_model.predict(x_test)

print('Ridge RMSE: ', mean_squared_error(y_test, ridge_preds, squared = False))
print('Ridge MAE: ', mean_absolute_error(y_test, ridge_preds))
```

```
Ridge RMSE:  0.01635486835867053
Ridge MAE:  0.013185163422887971
```

```
[31]: from sklearn.ensemble import RandomForestRegressor

# random forest regression
rf_model = RandomForestRegressor(n_estimators = 300,
                                  max_depth = 6,
                                  random_state = 42)
rf_model.fit(x_train, y_train)
rf_preds = rf_model.predict(x_test)
print('RF RMSE: ', mean_squared_error(y_test, rf_preds, squared = False))
print('RF MAE: ', mean_absolute_error(y_test, rf_preds))
```


```
RF RMSE:  0.02010146066469092
RF MAE:  0.01610909346292938
```

```
[32]: import matplotlib.pyplot as plt

# make sure y_test and predictions are aligned
test_weeks = model_df.loc[y_test.index, 'week']


# plot actual vs predicted change
plt.figure(figsize = (10, 5))
plt.plot(test_weeks, y_test.values, label = 'Actual')
plt.plot(test_weeks, rf_preds, label = 'RF Predicted')
plt.legend()
plt.title('3-Week Ahead ICU Occupancy Forecast')
plt.xticks(rotation = 45)
plt.tight_layout()
```

```
plt.show()
```



```
[33]: import pandas as pd
```

```
importances = pd.Series(rf_model.feature_importances_,  
                        index = x.columns).sort_values(ascending = False)  
  
importances.plot(kind = 'bar', figsize = (8,4))  
plt.title('Random Forest Feature Importance')  
plt.tight_layout()  
plt.show()
```



```
[34]: # update forecast to 7-weeks
forecast_horizon = 7

data['icu_occ_t_plus_7w'] = data['icu_occ'].shift(-7)
model_df = data.dropna(subset = features + ['icu_occ_t_plus_7w']).copy()
```

```
[35]: # select predictors
feature_candidates = ['icu_occ', 'FLU_ED_VISITS_per_100k',
                      'RSV_ED_VISITS_per_100k',
                      'FLU_ED_VISITS_per_100k_roll2',
                      'FLU_ED_VISITS_per_100k_roll3',
                      'RSV_ED_VISITS_per_100k_roll2',
                      'RSV_ED_VISITS_per_100k_roll3']
features = [f for f in feature_candidates if f in model_df.columns]
print('Using features: ', features)
```

Using features: ['icu_occ']

```
[36]: x = model_df[features].copy()
y = model_df['icu_occ_t_plus_3w']

# drop rows with missing predictors
mask = x.notna().all(axis = 1)
x = x[mask]
y = y[mask]

print('Final Modeling Rows: ', len(x))

# train/test split
split_idx = int(len(x) * 0.8)

x_train = x.iloc[:split_idx]
x_test = x.iloc[split_idx:]

y_train = y.iloc[:split_idx]
y_test = y.iloc[split_idx:]
```

Final Modeling Rows: 191

```
[37]: # build ridge model
ridge_model = Pipeline([('scaler', StandardScaler()),
                       ('ridge', Ridge(alpha = 1.0))])
ridge_model.fit(x_train, y_train)
ridge_preds = ridge_model.predict(x_test)

print('Ridge RMSE: ', mean_squared_error(y_test, ridge_preds, squared = False))
```

```
print('Ridge MAE: ', mean_absolute_error(y_test, ridge_preds))
```

Ridge RMSE: 0.015486797449980468

Ridge MAE: 0.012732300156640845

[38]: # build random forest model

```
rf_model = RandomForestRegressor(n_estimators = 300,
                                 max_depth = 6,
                                 random_state = 42)
rf_model.fit(x_train, y_train)
rf_preds = rf_model.predict(x_test)
print('RF RMSE: ', mean_squared_error(y_test, rf_preds, squared = False))
print('RF MAE: ', mean_absolute_error(y_test, rf_preds))
```

RF RMSE: 0.020534437222615023

RF MAE: 0.01641447692032386

[39]: data = data.sort_values('week').reset_index(drop=True)

```
# keep only rows where current ICU occupancy exists
data2 = data.dropna(subset = ['icu_occ']).copy()

FORECAST_HORIZON = 7
data2['icu_occ_t_plus_7w'] = data2['icu_occ'].shift(-FORECAST_HORIZON)
data2['icu_change_7w'] = data2['icu_occ_t_plus_7w'] - data2['icu_occ']

print(data2[['week', 'icu_occ', 'icu_occ_t_plus_7w', 'icu_change_7w']].head(12))
print('Rows with target available:', data2['icu_change_7w'].notna().sum())
```

	week	icu_occ	icu_occ_t_plus_7w	icu_change_7w
20	2020-07-07	0.084892	0.154956	0.070065
21	2020-07-14	0.105866	0.119519	0.013653
22	2020-07-21	0.137915	0.148674	0.010759
23	2020-07-28	0.134487	0.142428	0.007940
24	2020-08-04	0.133706	0.143664	0.009959
25	2020-08-11	0.128265	0.150275	0.022011
26	2020-08-18	0.136613	0.158488	0.021875
27	2020-08-25	0.154956	0.142436	-0.012520
28	2020-09-01	0.119519	0.156236	0.036717
29	2020-09-08	0.148674	0.150128	0.001454
30	2020-09-15	0.142428	0.130297	-0.012131
31	2020-09-22	0.143664	0.121798	-0.021866

Rows with target available: 191

[40]: # sort chronologically

```
data = data.sort_values('week').reset_index(drop=True)

# keep only rows where ICU occupancy exists
data2 = data.dropna(subset=['icu_occ']).copy()
```

```

print('Rows with valid ICU:', len(data2))

# create 7-week horizon change
FORECAST_HORIZON = 7

data2['icu_occ_t_plus_7w'] = data2['icu_occ'].shift(-FORECAST_HORIZON)
data2['icu_change_7w'] = data2['icu_occ_t_plus_7w'] - data2['icu_occ']

# drop rows without future target
model_df = data2.dropna(subset = ['icu_change_7w']).copy()

print('Model rows:', len(model_df))
print('Date range:',
      model_df['week'].min(),
      '→',
      model_df['week'].max())

```

Rows with valid ICU: 198
 Model rows: 191
 Date range: 2020-07-07 00:00:00 → 2024-02-27 00:00:00

```
[41]: # resplit
features = ['icu_occ', 'FLU_ED_VISITS_per_100k', 'RSV_ED_VISITS_per_100k']

features = [f for f in features if f in model_df.columns]

split_idx = int(len(model_df) * 0.8)

train_df = model_df.iloc[:split_idx]
test_df = model_df.iloc[split_idx:]

X_train = train_df[features]
y_train = train_df['icu_change_7w']

X_test = test_df[features]
y_test = test_df['icu_change_7w']
```

```
[42]: from sklearn.metrics import mean_squared_error, mean_absolute_error

# naive baseline: predict zero change
naive_preds = np.zeros(len(y_test))

naive_rmse = mean_squared_error(y_test, naive_preds, squared=False)
naive_mae = mean_absolute_error(y_test, naive_preds)

print('Naive RMSE:', naive_rmse)
print('Naive MAE:', naive_mae)
```

```
Naive RMSE: 0.024170987362290043  
Naive MAE: 0.02001658917467064
```


```
[43]: # build new ridge model  
ridge_model = Pipeline([('scaler', StandardScaler()),  
                      ('ridge', Ridge(alpha = 1.0))])  
  
ridge_model.fit(X_train, y_train)  
ridge_preds = ridge_model.predict(X_test)  
  
ridge_rmse = mean_squared_error(y_test, ridge_preds, squared = False)  
ridge_mae = mean_absolute_error(y_test, ridge_preds)  
  
print('Ridge RMSE:', ridge_rmse)  
print('Ridge MAE:', ridge_mae)
```

```
Ridge RMSE: 0.02490003351377585  
Ridge MAE: 0.020886145624497515
```

```
[44]: from sklearn.ensemble import RandomForestRegressor  
  
# build new random forest model  
rf_model = RandomForestRegressor(n_estimators = 300,  
                                 max_depth = 6,  
                                 random_state = 42)  
  
rf_model.fit(X_train, y_train)  
rf_preds = rf_model.predict(X_test)  
  
rf_rmse = mean_squared_error(y_test, rf_preds, squared=False)  
rf_mae = mean_absolute_error(y_test, rf_preds)  
  
print('RF RMSE:', rf_rmse)  
print('RF MAE:', rf_mae)
```

```
RF RMSE: 0.02845821680279954  
RF MAE: 0.021704250602392942
```

```
[45]: # plot actual vs predicted in 7-week ahead model  
plt.figure(figsize = (10, 5))  
plt.plot(test_df['week'], y_test.values, label = 'Actual Change')  
plt.plot(test_df['week'], ridge_preds, label = 'Ridge Predicted')  
plt.axhline(0, linestyle = '--', alpha = 0.5)  
plt.legend()  
plt.title('7-Week Ahead ICU Occupancy Change Forecast')  
plt.xticks(rotation = 45)  
plt.tight_layout()  
plt.show()
```



```
[46]: # identify ED-rate columns that actually exist
flu_col = next((c for c in model_df.columns if 'flu' in c.lower() and
                'per_100k' in c.lower()), None)
rsv_col = next((c for c in model_df.columns if 'rsv' in c.lower() and
                'per_100k' in c.lower()), None)

print('Detected flu column:', flu_col)
print('Detected rsv column:', rsv_col)

# add lagged ICU levels
model_df = model_df.sort_values('week').reset_index(drop = True)
model_df['icu_lag1'] = model_df['icu_occ'].shift(1)
model_df['icu_lag2'] = model_df['icu_occ'].shift(2)

# add lagged ED signals if columns exist
if flu_col is not None:
    model_df['flu_lag1'] = model_df[flu_col].shift(1)
    model_df['flu_lag2'] = model_df[flu_col].shift(2)

if rsv_col is not None:
    model_df['rsv_lag1'] = model_df[rsv_col].shift(1)
    model_df['rsv_lag2'] = model_df[rsv_col].shift(2)

# drop rows with new NaNs from lags/targets
model_df = model_df.dropna(subset = ['icu_change_7w', 'icu_occ', 'icu_lag1', 'icu_lag2']).copy()
if flu_col is not None:
    model_df = model_df.dropna(subset = ['flu_lag1', 'flu_lag2']).copy()
```

```

if rsv_col is not None:
    model_df = model_df.dropna(subset = ['rsv_lag1', 'rsv_lag2']).copy()

print('Rows after adding lags:', len(model_df))

```

Detected flu column: None
 Detected rsv column: None
 Rows after adding lags: 189

```
[47]: features = ['icu_occ', 'icu_lag1', 'icu_lag2']

if flu_col is not None:
    features += [flu_col, 'flu_lag1', 'flu_lag2']
if rsv_col is not None:
    features += [rsv_col, 'rsv_lag1', 'rsv_lag2']

print('Features used:', features)
```

Features used: ['icu_occ', 'icu_lag1', 'icu_lag2']

```
[48]: flu_col = 'percent_visits_smoothed_influenza'
rsv_col = 'percent_visits_smoothed_rsv'

print('Using flu column:', flu_col)
print('Using rsv column:', rsv_col)
```

Using flu column: percent_visits_smoothed_influenza
 Using rsv column: percent_visits_smoothed_rsv

```
[49]: model_df = model_df.sort_values('week').reset_index(drop=True)

# ICU lags
model_df['icu_lag1'] = model_df['icu_occ'].shift(1)
model_df['icu_lag2'] = model_df['icu_occ'].shift(2)

# ED lags
model_df['flu_lag1'] = model_df[flu_col].shift(1)
model_df['flu_lag2'] = model_df[flu_col].shift(2)

model_df['rsv_lag1'] = model_df[rsv_col].shift(1)
model_df['rsv_lag2'] = model_df[rsv_col].shift(2)

# drop rows with lag-induced NaNs
model_df = model_df.dropna().copy()

print('Rows after adding lags:', len(model_df))
```

Rows after adding lags: 69

```
[50]: features = ['icu_occ', 'icu_lag1', 'icu_lag2', flu_col,
                 'flu_lag1', 'flu_lag2', rsv_col, 'rsv_lag1', 'rsv_lag2']

print('Feature count:', len(features))
```

Feature count: 9

```
[51]: split_idx = int(len(model_df) * 0.8)

train_df = model_df.iloc[:split_idx]
test_df = model_df.iloc[split_idx:]

X_train = train_df[features]
y_train = train_df['icu_change_7w']

X_test = test_df[features]
y_test = test_df['icu_change_7w']
```

```
[52]: naive_preds = np.zeros(len(y_test))

print('Naive RMSE:', mean_squared_error(y_test, naive_preds, squared=False))
print('Naive MAE:', mean_absolute_error(y_test, naive_preds))
```

Naive RMSE: 0.023513781015979705

Naive MAE: 0.019637858420002756

```
[53]: ridge_model = Pipeline([('scaler', StandardScaler()),
                            ('ridge', Ridge(alpha = 1.0))])

ridge_model.fit(X_train, y_train)
ridge_preds = ridge_model.predict(X_test)

print('Ridge RMSE:', mean_squared_error(y_test, ridge_preds, squared = False))
print('Ridge MAE:', mean_absolute_error(y_test, ridge_preds))
```

Ridge RMSE: 0.02478472755135359

Ridge MAE: 0.020009696341092075

```
[54]: rf_model = RandomForestRegressor(n_estimators = 300,
                                      max_depth = 6,
                                      random_state = 42)

rf_model.fit(X_train, y_train)
rf_preds = rf_model.predict(X_test)

print('RF RMSE:', mean_squared_error(y_test, rf_preds, squared = False))
print('RF MAE:', mean_absolute_error(y_test, rf_preds))
```

RF RMSE: 0.021597451258293918


RF MAE: 0.017485906630914816

```
[55]: importances = pd.Series(rf_model.feature_importances_,
                               index = features).sort_values(ascending = False)

print(importances)

plt.figure(figsize = (8, 4))
importances.plot(kind = 'bar')
plt.title('Random Forest Feature Importance')
plt.tight_layout()
plt.show()
```

```
icu_lag2          0.278436
icu_occ           0.249924
percent_visits_smoothed_rsv 0.110652
percent_visits_smoothed_influenza 0.097884
icu_lag1          0.097785
rsv_lag1          0.056424
flu_lag2          0.038021
rsv_lag2          0.037510
flu_lag1          0.033363
dtype: float64
```



```
[56]: # Create directional target
model_df['increase_7w'] = (model_df['icu_change_7w'] > 0).astype(int)
print(model_df['increase_7w'].value_counts())
```

```
increase_7w
1    42
```

```
0      27  
Name: count, dtype: int64
```

```
[57]: features = ['icu_occ', 'icu_lag1', 'icu_lag2',  
    ↴'percent_visits_smoothed_influenza',  
        'flu_lag1', 'flu_lag2', 'percent_visits_smoothed_rsv', 'rsv_lag1',  
    ↴'rsv_lag2']  
  
model_df = model_df.sort_values('week').reset_index(drop = True)  
  
split_idx = int(len(model_df) * 0.8)  
  
train_df = model_df.iloc[:split_idx]  
test_df  = model_df.iloc[split_idx:]  
  
X_train = train_df[features]  
y_train = train_df['increase_7w']  
  
X_test  = test_df[features]  
y_test  = test_df['increase_7w']
```

```
[58]: from sklearn.metrics import accuracy_score  
  
majority_class = y_train.mode()[0]  
baseline_preds = np.full(len(y_test), majority_class)  
  
print('Baseline Accuracy:', accuracy_score(y_test, baseline_preds))
```

```
Baseline Accuracy: 0.7142857142857143
```

```
[59]: from sklearn.ensemble import RandomForestClassifier  
from sklearn.metrics import (accuracy_score, roc_auc_score,  
    ↴classification_report, confusion_matrix)  
  
rf_clf = RandomForestClassifier(n_estimators = 400,  
                                max_depth = 6,  
                                random_state = 42)  
  
rf_clf.fit(X_train, y_train)  
  
rf_probs = rf_clf.predict_proba(X_test)[:, 1]  
rf_preds = (rf_probs > 0.5).astype(int)  
  
print('RF Accuracy:', accuracy_score(y_test, rf_preds))  
print('RF ROC-AUC:', roc_auc_score(y_test, rf_probs))  
  
print('\nConfusion Matrix:')print(confusion_matrix(y_test, rf_preds))
```

```
print('\nClassification Report:')

print(classification_report(y_test, rf_preds))
```

RF Accuracy: 0.8571428571428571

RF ROC-AUC: 1.0

Confusion Matrix:

```
[[4 0]
 [2 8]]
```


Classification Report:

	precision	recall	f1-score	support
0	0.67	1.00	0.80	4
1	1.00	0.80	0.89	10
accuracy			0.86	14
macro avg	0.83	0.90	0.84	14
weighted avg	0.90	0.86	0.86	14

```
[60]: from sklearn.metrics import roc_curve

fpr, tpr, _ = roc_curve(y_test, rf_probs)

plt.figure(figsize = (6, 6))
plt.plot(fpr, tpr, label = f'RF (AUC = {roc_auc_score(y_test, rf_probs):.3f})')
plt.plot([0,1], [0,1], linestyle = '--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve - 7-Week ICU Increase Prediction')
plt.legend()
plt.tight_layout()
plt.show()
```




```
[61]: clf_importance = pd.Series(rf_clf.feature_importances_, index=features)
      .sort_values(ascending = False)

print(clf_importance)

clf_importance.plot(kind = 'bar', figsize = (8, 4))
plt.title('RF Classification Feature Importance')
plt.tight_layout()
plt.show()
```

icu_occ	0.243025
icu_lag2	0.185907
icu_lag1	0.167263

```
percent_visits_smoothed_influenza      0.102939
flu_lag1                                0.070255
flu_lag2                                0.066988
percent_visits_smoothed_rsv             0.064068
rsv_lag1                                 0.059048
rsv_lag2                                 0.040506
dtype: float64
```



```
[ ]:
```